

Exploratory Data Analysis (EDA)

Univariate Analysis

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Load the dataset (replace with your file path)
data = pd.read_csv('/content/heart disease prediction.csv')

# Display basic information about the dataset
print(data.head())
print(data.info())
print(data.describe())

# Univariate Analysis
# Distribution of the target variable
plt.figure(figsize=(6, 4))
sns.countplot(x='HeartDisease', data=data)
plt.title('Distribution of Heart Disease')
plt.show()

# Distribution of numerical features
data.hist(bins=30, figsize=(15, 10))
plt.show()
```

```

0  40  M      ATA      140      289      0      Normal      172
1  49  F      NAP      160      180      0      Normal      156
2  37  M      ATA      130      283      0      ST          98
3  48  F      ASY      138      214      0      Normal      108
4  54  M      NAP      150      195      0      Normal      122

```

```

ExerciseAngina  Oldpeak  ST_Slope  HeartDisease
0              N      0.0      Up          0
1              N      1.0      Flat        1
2              N      0.0      Up          0
3              Y      1.5      Flat        1
4              N      0.0      Up          0

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 918 entries, 0 to 917
```

```
Data columns (total 12 columns):
```

```

#   Column              Non-Null Count  Dtype
---  -
0   Age                918 non-null    int64
1   Sex                918 non-null    object
2   ChestPainType       918 non-null    object
3   RestingBP           918 non-null    int64
4   Cholesterol          918 non-null    int64
5   FastingBS           918 non-null    int64
6   RestingECG          918 non-null    object
7   MaxHR               918 non-null    int64
8   ExerciseAngina      918 non-null    object
9   Oldpeak             918 non-null    float64
10  ST_Slope            918 non-null    object
11  HeartDisease         918 non-null    int64

```

```
dtypes: float64(1), int64(6), object(5)
```

```
memory usage: 86.2+ KB
```

```
None
```

```

      Age  RestingBP  Cholesterol  FastingBS  MaxHR  \
count  918.000000  918.000000  918.000000  918.000000  918.000000
mean    53.510893  132.396514  198.799564    0.233115  136.809368
std     9.432617   18.514154   109.384145    0.423046   25.460334
min     28.000000    0.000000    0.000000    0.000000   60.000000
25%     47.000000   120.000000   173.250000    0.000000  120.000000
50%     54.000000   130.000000   223.000000    0.000000  138.000000
75%     60.000000   140.000000   267.000000    0.000000  156.000000
max     77.000000   200.000000   603.000000    1.000000  202.000000


```

```

      Oldpeak  HeartDisease
count  918.000000  918.000000
mean     0.887364    0.553377
std     1.066570    0.497414
min     -2.600000    0.000000
25%      0.000000    0.000000
50%      0.600000    1.000000
75%      1.500000    1.000000
max      6.200000    1.000000

```

Distribution of Heart Disease



Bivariate Analysis

```
# Correlation heatmap
```

```
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(data.select_dtypes(include=np.number).corr(), annot=True, cmap='coolwarm', vmin=-1, vmax=1) # Select only numerical columns f
```

```
plt.title('Correlation Heatmap')
```

```
plt.show()
```


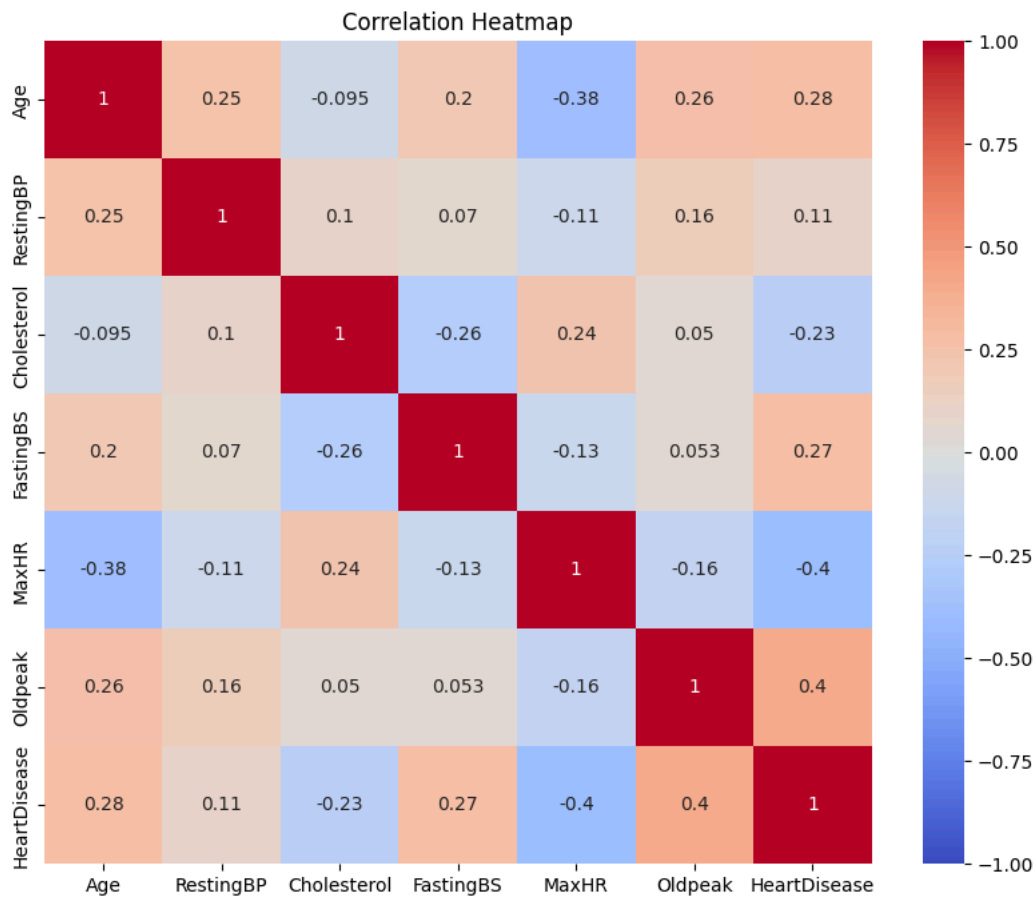
```
# Boxplot for numerical features by target
```

```
plt.figure(figsize=(12, 6))
```

```
sns.boxplot(x='HeartDisease', y='Age', data=data)
```

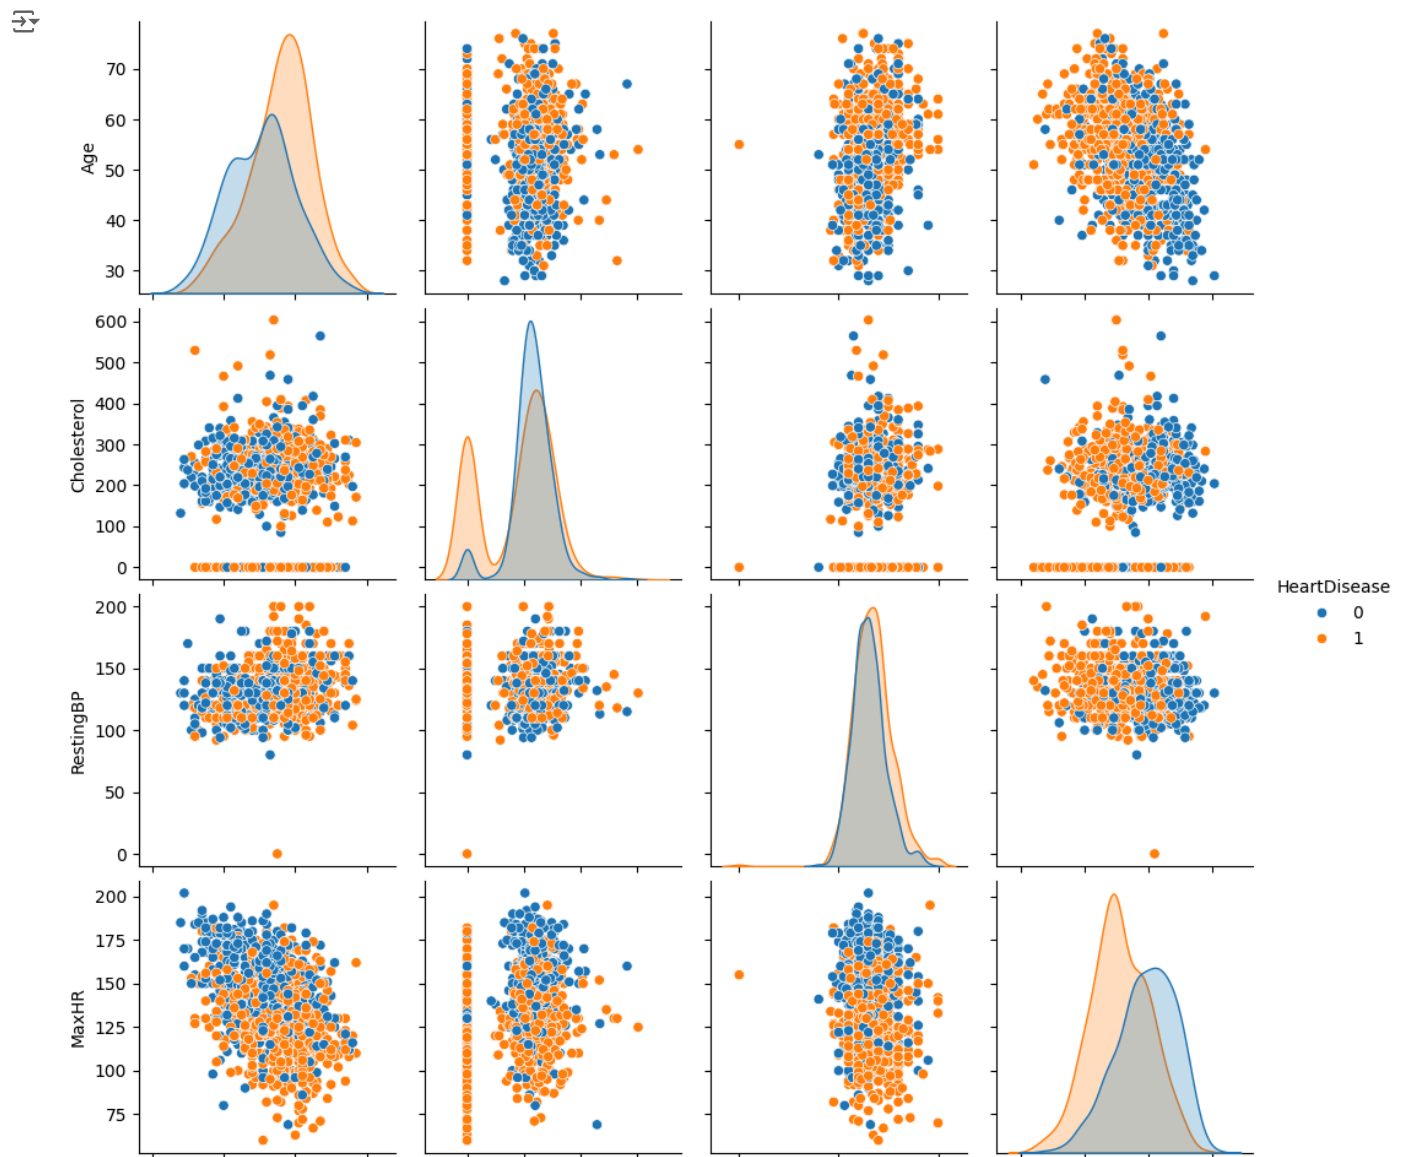
```
plt.title('Age vs Heart Disease')
```

```
plt.show()
```



Multivariate Analysis

```
# Pairplot for selected features
sns.pairplot(data, hue='HeartDisease', vars=['Age', 'Cholesterol', 'RestingBP', 'MaxHR'])
plt.show()
```



Data Preprocessing

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score
import joblib

# Preprocessing
# Handle missing values and encode categorical variables
features = data.drop(columns=['HeartDisease'])
target = data['HeartDisease']

# Identify categorical and numerical columns
categorical_features = features.select_dtypes(include=['object']).columns
numerical_features = features.select_dtypes(exclude=['object']).columns

# Define preprocessing for numerical and categorical data
numerical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_transformer, numerical_features),
        ('cat', categorical_transformer, categorical_features)])

# Split data
```

```
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Define model
model = Pipeline(steps=[('preprocessor', preprocessor),
                        ('classifier', RandomForestClassifier(random_state=42))])

# Train model
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
print("Train Accuracy:", accuracy_score(y_train, model.predict(X_train)))
print("Test Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
↗ Train Accuracy: 1.0
Test Accuracy: 0.8804347826086957
```

	precision	recall	f1-score	support
0	0.85	0.87	0.86	77
1	0.90	0.89	0.90	107